

Lecture 1: Introduction

Class Grades Break Down:

- Quizzes/Class participation – 10%
- 3-4 programming assignments -45%
 - Every 3-4 weeks
 - weights may be uneven
 - Learn PHP and HTML by yourself
 - 25% per week deduction/maximum two weeks
- Three exams (45%)
 - 15% each
 - in-class, closed-book exams
 - No final exam
 - No make-ups

Database Management System (DBMS):

- A DBMS is a powerful tool for **creating** and **managing** large amounts of data *efficiently* and allowing it to persist over long periods of time *safely*.

What are the problems with file systems? (Reason behind why we converted to DBMS)

- Do not support efficient access to data items
- Do not have a query language for data items
- Do not support concurrent control

Relational Database Systems - RDBMS:

Relations = Data organized as **tables**

Queries could be expressed in high level language:

- No need for handling complex storage structure directly
- Greatly increased the efficiency of database programmers

	<p>Transaction Processing</p> <p>A transaction is a group of database operations with certain properties.</p> <ul style="list-style-type: none"> - E.g., bank transfer (withdraw + deposit) <p>The ACID properties of transactions</p> <ul style="list-style-type: none"> - A (Atomicity): all-or-nothing execution of transaction - C (Consistency): consistency constraints can not be violated - I (Isolation): each transaction must appear to be executed as if no other transaction is executing at the same time - D (Durability): the effect on the database of a transaction must not be lost once the transaction has completed <p>The transaction processor performs the following task to ensure ACID properties:</p> <ul style="list-style-type: none"> - Logging - Concurrency Control - Deadlock Resolution <p>IMPORTANT QUERY COMMANDS:</p> <p>SELECT * FROM [name of table];</p> <ul style="list-style-type: none"> - Displays all items in the table
<p>Lecture 2&3: The Relational Model of Data</p>	<p>What is a Data Model?</p> <p>A data model is a notation for describing data or info. The descriptions include:</p> <ol style="list-style-type: none"> 1. <i>Structure</i> of the data <ol style="list-style-type: none"> a. Relational Model = Tables b. Semi-structured model = tree/graphs 2. <i>Operations</i> on Data <ol style="list-style-type: none"> a. A limited set of queries & modifications allowed

- b. The limitation allows us to describe database operations at very high levels → efficient implementations in DBMS
- c. *Constraints*: limit on what data could be

Basics of Relational model:

- *Attributes*: Columns of a relation
- *Schema*: The name of a relation and the set of attributes for a relation
 - Example: Movies(title, year, length, genre)
 - The attributes are a SET, not a list
- A database consists of 1 or more relations.
- Database schema = a set of all relation schemas in the database
- Tuples: the rows of a relation (excluding the header)
 - Ex. (Star Wars, 1977, 120, SciFi)
 - It is just one row of data, a tuple has one component for each attribute of the relation
- Domains
 - Each component of each tuple must be atomic:
 - Must be a datatype (ex. Integer, string)
 - CAN NOT be a structure ex. Array, list
 - Each attribute of a relation is a domain(a particular elementary data type)

Movies (title:string, year:integer, length:integer, genre:string)

(Ask if lecture 2 slide 8 is correct)

Keys of Relations *

- A set of attributes that form a key for a relation if we do not allow 2 tuples in a relation instance to have the same values in all the attributes of the key.
- Tuples can not have the same value in all key attributes

Two ways of declaring keys by adding one of the following keywords after an attribute when declaring/altering a table

- **PRIMARY KEY** (the explanation of PRIMARY is deferred)
- **UNIQUE**

PRIMARY KEY vs. UNIQUE

1. There can be only one PRIMARY KEY for a relation, but several UNIQUE attributes.
2. No attribute of a PRIMARY KEY can ever be NULL in any tuple. But attributes declared as UNIQUE may have NULL's, and there may be several tuples with NULL.

SQL - Structured Query Language

- Principle language used to describe and manipulate relational databases.

Relations in SQL: three types

- Tables
- Views
- Temporary tables

Data Types

- Character strings
 - CHAR(n): a fixed-length string of up to n characters
 - VARCHAR(n): a string up to n characters
 - The difference is implementation dependent
 - Generally VARCHAR(n) is more space efficient
 - INT/INTEGER:
 - typical integer values.
 - Variants: SHORTINT (smallint for mysql/postgresql)
 - Float pointing numbers
 - **FLOAT/REAL (DOUBLE)**
 - **DECIMAL(n,d)**: values consist of n decimal digits with the decimal point assumed to be the d position from the right. → **NUMERIC (n,d)**
 - Dates and Time
 - DATE '1948-05-14'
 - TIME '15:00:02.5'
 - Bit strings/boolean

Simple table declarations:

CREATE TABLE ...

Example:

```
CREATE TABLE Movies
(
    title          CHAR(100),
    year           INT,
    genre          CHAR(10),
    studioName     CHAR(30),
    producerC#     INT
);
```

Modifying Relation Schemas:

- How to drop a table in sql:
 - **DROP TABLE [name]**
- Alter table
 - **ALTER TABLE [name] ADD ..**
 - **ALTER TABLE [name] DROP ..**

Examples:

DROP TABLE Movies;

ALTER TABLE MovieStar ADD phone CHAR(16);

ALTER TABLE MovieStar DROP birthdate;

Assigning Default Values

- When tuples are created or modified, sometimes we do not have values for all components → NULL are assigned to the missing values
- We can change NULL to something more appropriate by using the **DEFAULT** keyword when declaring or modifying an attribute

EXAMPLES:

1. gender CHAR(1) DEFAULT '?'
2. birthday DATE DEFAULT DATE '0000-00-00'
3. ALTER TABLE MovieStar ADD phone CHAR(16)

DEFAULT 'unlisted'

An Algebraic Query Language – Relational Algebra

What is Relational Algebra?

- An algebra whose operands are relations or variables that represent relations.
- Operators are designed to do the most common things that we need to do with relations in a database.
- The result is an algebra that can be used as a *query language* for relations.

Core Relational Algebra:

- Selection: picking certain rows
- Projection: picking certain columns
- Products & Join: compositions of relations
- Renaming of relations and attributes.
- Union, intersection and difference

Selection:

$$R1 := \sigma_C(R2)$$

- C is a condition (as in “if” statements) that refers to attributes of $R2$.
- $R1$ is all those tuples of $R2$ that satisfy C .

Look @Slide3.9

Projection:

$$R1 := \pi_L(R2)$$

- L is a list of attributes from the schema of $R2$.
- $R1$ is constructed by looking at each tuple of $R2$, extracting the attributes on list L , in the order specified, and creating from those components of a tuple for $R1$.
- Eliminate duplicate tuples, if any.

Look @Slide3.11

Extended Projection:

- Using the same π_L operator, we allow the list L to

contain arbitrary expressions involving attributes:

1. Arithmetic on attributes, e.g., $A+B \rightarrow C$.
2. Duplicate occurrences of the same attribute.

Look @Slide3.13

Product:

$R3 := R1 \times R2$

- Pair each tuple $t1$ of $R1$ with each tuple $t2$ of $R2$.
- Concatenation of $t1$ and $t2$ is a tuple of $R3$.
- Schema of $R3$ is the attributes of $R1$ and then $R2$, in the order.
- But beware attribute A of the same name in $R1$ and $R2$: use $R1.A$ and $R2.A$.

Natural Join:

Denoted $R3 := R1 \bowtie R2$.

- A *natural* join connects two relations by:
 - Equating attributes of the same name
 - Projecting out one copy of each pair of equated attributes.

Theta Join:

- $R3 := R1 \bowtie_C R2$
 - Take the product $R1 \times R2$.
 - Then apply σ_C to the result.
- As for σ , C can be any boolean-valued condition.

Building Complex Expressions

Three notations, just as in arithmetic:

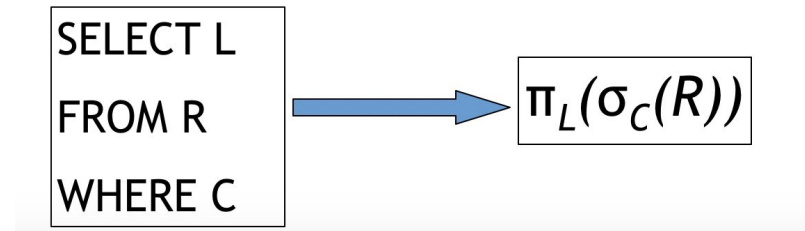
1. Expressions with several operators.

$$R' = \pi_L(\sigma_C(R))$$

2. Sequences of assignment statements.
3. Expression trees.

Select-From-Where Statements:

SELECT desired attributes
FROM one or more tables
WHERE condition about tuples of the table



* In **SELECT** clauses:

* in the **SELECT** clause stands for “all attributes of this relation.”

Example: Using Beers(name, manf):

```
SELECT *  
FROM Beers  
WHERE manf = 'Anheuser-Busch';
```

Renaming Attributes

If you want the result to have different attribute names, use “*AS <new name>*” to rename an attribute.

Example: Using Beers(name, manf):

```
SELECT name AS beer, manf  
FROM Beers  
WHERE manf = 'Anheuser-Busch'
```

Expressions in **SELECT** Clauses:

Any expression that makes sense can appear as an element of a **SELECT** clause.

Example: Using Sells(bar, beer, price):

```
SELECT bar, beer,  
price*114 AS priceInYen  
FROM Sells;
```


@Slide3.17

Complex Conditions in WHERE Clause:

- Boolean operators AND, OR, NOT.
- Comparisons =, <>, <, >, <=, >=.

Using Sells(bar, beer, price), find the price Joe's Bar charges for Bud:

```
SELECT price
FROM Sells
WHERE bar = 'Joe's Bar' AND
      beer = 'Bud';
```

Using Drinkers(name, addr, phone) find the drinkers with exchange 555:

```
SELECT name
FROM Drinkers
WHERE phone LIKE '%555-____';
```

NULL Values:

- *Missing value* : e.g., we know Joe's Bar has some address, but we don't know what it is.
- *Inapplicable* : e.g., the value of attribute spouse for an unmarried person.

Multi-relation Queries:

- Dealing with more than one relation aka table
- You can distinguish attributes of the same name by "<relation aka table name>.<attribute>"

Using relations Likes(drinker, beer) and Frequents(drinker, bar), find the beers liked by at least one person who frequents Joe's Bar:

```
SELECT beer  
  
FROM Likes, Frequents  
WHERE bar = 'Joe's' AND  
Frequents.drinker = Likes.drinker;
```

How to do multi-relation Queries:

1. Start with the **product of all the relations** in the FROM clause.
2. Apply the **selection condition** from the WHERE clause.
3. Project onto the **list of attributes and expressions** in the SELECT clause.

Slide@5.7

SELF-JOIN:

From Beers(name, manf), find all pairs of beers by the same manufacturer.

- Do not produce pairs like (Bud, Bud).
- Produce pairs in alphabetic order, e.g. (Bud, Miller), not (Miller, Bud).

```
SELECT b1.name, b2.name  
FROM Beers b1, Beers b2  
WHERE b1.manf = b2.manf AND  
b1.name < b2.name;
```

Subqueries:

- A parenthesized SELECT-FROM-WHERE statement is a (*subquery*)

Subquery of the previous question:

Find the beers liked by someone who frequents Joe's Bar.

```
SELECT beer
FROM LIKES, (SELECT drinker FROM Frequents WHERE
bar = "Joe's")JD
WHERE Likes.drinker = JD.drinker;
```

Using Sells(bar, beer, price), find the bars that serve Miller for the same price Joe's bar charges for Bud.

```
SELECT bar
FROM Sells
WHERE beer = 'Miller' AND
      price = (SELECT price
FROM Sells
WHERE bar = 'Joe's Bar'
AND beer = 'Bud');
```

The IN & NOT IN Operator:

<tuple> IN (subquery)

- is true if and only if the tuple is a member of the relation produced by the subquery.
- Opposite: <tuple> NOT IN (<subquery>).

Examples: Comparison

```
SELECT a
FROM R, S
WHERE R.b = S.b;
```

```
SELECT a
FROM R
WHERE b IN (SELECT b FROM S);
```

@Slide5.21-26

Lecture 6:

Relational Algebra on Bags:

- A bag or multiset is like a set, but an element may appear more than once.
- Example:
 - {1,2,3,1} is a bag
 - {1,2,3} is also a bag that happens to be a set

Why Bags?

- The bag language is the most important query language for relational databases.

Operations on Bags:

- Selection
- Projection (allows duplicated in bag operand)
- Products and Joins

Go over Slides to get info on UNION, INTERSECTION, DISTINCT, ALL...

δ = eliminate duplicates from bags.

τ = sort tuples.

γ = grouping and aggregation.

Outerjoin : avoids “dangling tuples” = tuples that do not join with anything

@Slide7.7

Aggregation Operators:

- The most important examples: SUM, AVG, COUNT(aka len), MIN, and MAX.

@Slide7.12

@Slide7.14

- Product:

R CROSS JOIN S;

- Natural join:

R NATURAL JOIN S;

- Example:

Likes NATURAL JOIN Sells;

- Relations can be (parenthesized) subqueries, as well.
- We may follow a SELECT-FROM-WHERE expression by GROUP BY and a list of attributes.
- HAVING <condition> may follow a GROUP BY clause.
 - If so, the condition applies to each group, and groups not satisfying the condition are eliminated.
- A *modification* command does not return a result (as a query does), but changes the database in some way.
- Three kinds of modifications:
 1. *Insert* a tuple or tuples.
 2. *Delete* a tuple or tuples.
 3. *Update* the value(s) of an existing tuple or tuples.
 4. To insert a single tuple:

INSERT INTO <relation>

VALUES (<list of values>);

- To delete tuples satisfying a condition from some relation:

DELETE FROM <relation>

WHERE <condition>;

Updates

To change certain attributes in certain tuples of a relation:

UPDATE <relation>

SET <list of attribute assignments>

WHERE <condition on tuples>;